

Comparative API Complexity Analysis of Two Platforms for Networked Multiplayer Games

Toni Alatalo

Department of Computer Science and Engineering Center for Internet Excellence (CIE)
University of Oulu,
Email: toni@playsign.net

Erno Kuusela

University of Oulu,
Email: erno@playsign.net

Rauli Puuperä

Playsign Ltd.
Oulu, Finland
Email: therauli@playsign.net

Abstract—In this paper we make a two-fold contribution towards understanding how gaming platform APIs succeed in easing the development of networked multiplayer games. First, we present an open source tool for automatic quantitative assessment of the complexity of a software API by analysing a source code based on the API with Sneed’s Object-Point (OP) method. Second, we apply the tool to compare the API complexity of the realXtend Tundra SDK to that of the Union platform using pre-existing implementations of the Pong game atop the two platforms. Our findings show that the tool successfully quantifies the different amounts of complexity that the game developer has to cope with on these two alternate platforms.

I. INTRODUCTION

Programming distributed networked applications is generally more complex than developing standalone software. Developer typically needs to deal with events, conflicts and error conditions originating from other parts of the distributed system as well as the local user. This is emphasized in multiuser real-time systems when compared to the relatively leisurely request-response interaction patterns of most client-server applications. Developing a massively multiplayer online game (MMOG) has been estimated to typically take two to three times longer than creating and launching a single-player game [1].

Higher level software abstractions are a common way to attempt to ease application development. Networking libraries simplify connection management and messaging. Distributed object systems automate remote calls and data synchronization. For an application developer, these resources are provided as a set of abstractions forming the application development interface (API). It has been noted how making good APIs is hard - and that creating a bad API is easy [2]. Even a small quirk in an API can accumulate to substantial problems in larger bodies of application source code. API design has a significant impact on software quality, and increased API complexity is associated with increased software failure rates [3].

The primary objective of this study is to evaluate how the API of the open source realXtend Tundra SDK (Tundra from now on), recently introduced by Alatalo [4], succeeds in hiding the complexity of developing networked multiplayer games. Tundra strives to adopt the best design practices from game engine literature, notably aggregation using an

abstract entity-component model. The objective of the entity-component model and the whole Tundra platform is to make development of networked multiplayer games easy and productive. Specifically in terms of networking, Tundra features attribute autosynchronization, a simple form of transparent remote procedure calls (entity actions) and efficient customized movement messages with inter- and extrapolation logic (dead reckoning).

Rigorous quantitative assessment of a gaming platform in terms of ease and productivity of development is very challenging. Hsiao and Yuan [1] proposed four essential “ease of” requirements for a MMOG middleware: ease of development, deployment, maintenance and change. However, they noted the difficulty of quantitatively measuring for example the ease of development or change, and focused only on platform scalability in their evaluation. In this paper we quantify the “ease of” dimension as an inverse function of the complexity of the software API used for game development. Inspired by recent progress in quantitative techniques for software complexity analysis [5], [3], we first propose a tool for automated quantitative assessment of the complexity of a software API based on automatic analysis of the source code using the API with Sneed’s Object-Point (OP) method. Then we use the tool to compare Tundra’s API complexity to that of the Union platform with pre-existing implementations of the Pong game on the two platforms, i.e. Pong serves as the minimal reference of a networked multiplayer game in our study. Our data shows that the tool indeed successfully quantifies the different amounts of complexity that the game developer has to manage on these two alternate platforms.

This paper is organized as follows. After briefly reviewing related work on API complexity analysis we present our tool. Then we report our case study on using the tool to compare the API complexity of the two platforms in the development of the Pong game. We conclude the paper with a discussion on various aspects of the study.

II. RELATED WORK ON API COMPLEXITY ANALYSIS

A wide range of qualitative and quantitative approaches have been proposed for assessing the complexity of software APIs. The API usability research has adapted traditional qualitative usability evaluation methods of the HCI field such as thinking aloud, heuristic evaluation and cognitive

walkthroughs into evaluating APIs [6]. Employed metrics include the completion times of predefined tasks, for example. In comparison to task-based usability evaluation of GUIs, API evaluation based on human observations is challenging. When developing a small application can take even weeks, fitting valid evaluation tasks into an observation session lasting typically 1-2 hours is difficult [7]. The recently introduced peer review [8] and the concept maps [7] methods address this problem by involving real world usage of an API over a long period of time. However, they are still considerably laborious and yield only qualitative findings.

Recently, statistical (quantitative) methods have been proposed for analysing the complexity of software APIs. Cataldo and de Sousa [3] studied two large corporate software projects and nine open source projects and found a link between the API complexity and the failure proneness of the software quantified by the number of bug reports from the field. They quantified the complexity of an API by simply calculating the number of public methods and attributes.

de Souza and Bentolila [9] introduced the Metrix tool that evaluates an API by calculating Bandi's software metrics for interface size, interaction level and operation argument complexity from the API specification. After analysing eleven different APIs with the Metrix tool, de Souza and Bentolila concluded that calculating such simple metrics directly from the API specification can produce misleading results. A more complete API may appear more complex, even if it provides good abstractions that allow completing a particular task at hand with only a small subset of the API.

An alternative to evaluating the complexity of an API with simple metrics calculated from the API specification is to assess the complexity of programs developed using the API. Sobernig et al. [5] proposed exactly this in their comparative study of the complexity of four different API designs. To characterize the complexity of a particular API, they applied Sneed's [10] Object-Point (OP) analysis to a software realized atop the API. The key is to apply a surrogate measurement so that a program developed using an API is analyzed, not the API specification itself.

III. TOOL FOR AUTOMATED ANALYSIS OF API COMPLEXITY

Our objective was to develop a tool that facilitates automated and quantitative evaluation of the complexity of an API from an existing body of source code developed using the API. Such a tool would come with several advantages. First, real world data, i.e. source codes of existing applications, could be utilized. Second, in comparison to manual methods the analysis would be quick with immediate feedback and would require only a limited amount of human labor. Third, longitudinal studies of API development would be straightforward to conduct by analysing successive software versions. A fully automated analysis could be embedded into the continuous integration of a software bundle, to characterize the evolution of its complexity over time.

Following the work of Sobernig et al. [5], we base our tool on Sneed's Object-Point (OP) method [10]. However, in contrast to their manual data collection, our tool automatically extracts the data needed by the OP method from a program's source code. The OP method uses intermediate UML models as the data to compare programs in different languages. Importantly, the OP method allows direct tracking between indicator values and program structures, which is elementary in evaluating API designs. For example, if many codebases get a high proportion of their complexity value due to a specific part of an API, it can then be examined qualitatively. So called API hotspots and coldspots have been previously automatically mined from source code [11]. There, however, the specific parts of an API are not analyzed as sources of complexity, but simply to identify how much they are used. Their source code mining is similar to our tool that also needs to identify which functions are called and how often to employ the OP method.

A. Sneed's Object-Point (OP) method

Although the OP method was originally developed for deriving early work estimates from UML design diagrams, recently it has been applied to analysing the complexity and cost of existing software implementations. While the early COCOMO software cost models used simply program size (LOC, lines of code) to estimate development effort, later the more versatile Function-Point, the Data-Point and finally the Object-Point methods have emerged to incorporate functionality and other program properties into the cost estimate [12].

Sobernig *et al.* illustrated the relative robustness of the OP method to the simple LoC (lines of code) metric on two software implementations. The first software had only 48 LoC but resulted in 356.34 OP. The second software had 144 LoC but only 266.76 OP. Their reasoning was: "*an API user is only exposed to an API feature chunk of low structural complexity*", as the chunk's "*size is limited in terms of participating classes and the smallest number of operations per class*" and it "*shows a relatively weak connectedness of classes, resulting from the small number of associations and generalizations between the classes*". This reasoning is of utmost relevance to our objective of easing the development of networked game development with good API designs. We pursue a limited set of powerful abstractions with clear interactions that a game developer could easily learn and grow to master. Not all source code lines are equal - a poorly designed API makes it a struggle to get even a few operations working if the developer has to strive for functionality scattered around in an incoherent way.

The Object-Points, as applied here, are a sum of two parts: Class Points (CP) and Message Points (MP).

Class Points (CP) (Eq. 1-4) are calculated from the static class structure: the class count and sums of attribute, operation and relation counts. Weights are employed to correct the values for the overall calculation. Class inheritance is taken into account by calculating novelty weights for specializing classes.

Message Points (MP) (Eq. 5-8) are defined by the set of operations (functions/methods) *actually used* in the software.

First, the number of operations is recorded. Then, the parameter count for each called operation is collected. Also, the source and target counts of the operation calls are established. Again, novelty weights are used to compensate for repeated occurrences due to subclassing.

$$CP = \left(W_C |C| + \sum_{c \in C} |A_c| + W_{R_c} \sum_{c \in C} |R_c| + W_{O_c} \sum_{c \in C} |O_c| \right) \overline{N_C}, \quad (1)$$

where

$$\overline{N_C} = \frac{\sum_{c \in C} N_c}{|C|}, \text{ and} \quad (2)$$

$$N_c = \begin{cases} 1, & \text{if class is novel} \\ 0.5, & \text{if class has a super class} \end{cases} \quad (3)$$

$$MP = \left(W_{O_M} |O_M| + \sum_{o \in O_M} |P_o| + W_{S_o} \sum_{o \in O_M} |S_o| + W_{T_o} \sum_{o \in O_M} |T_o| \right) \overline{N_{O_M}}, \quad (4)$$

where

$$\overline{N_{O_M}} = \frac{\sum_{o \in O_M} N_o}{|O_M|}, \text{ and} \quad (5)$$

$$N_o = \begin{cases} 1, & \text{if operation is novel} \\ 0.5, & \text{if operation is provided by super class} \end{cases} \quad (6)$$

$C, |C| \dots$ Set of classes, Class count

$|A_c| \dots$ Attribute count per class

$|O_c| \dots$ Operation count --"

$|R_c| \dots$ Relation count --"

$N_c \dots$ Novelty weight of class c

$\overline{N_C} \dots$ Avg. class novelty

$|O_M| \dots$ Set of called operations

$|O_M| \dots$ Called operation counts

$|P_o| \dots$ Parameter count of operation o

$|S_o| \dots$ Source count --"

$|T_o| \dots$ Target count --"

$N_o \dots$ Novelty weight --"

The weights are adopted directly from the earlier usage of OP for API complexity analysis, which further uses the standard Data Points analysis values by Sneed. Please see [5] for a detailed description of the calculation of the CP and MP values.

B. Extracting Object-Point data from source code

To obtain the static class data for the Class Points (CP), we utilize existing source code parsing and annotation systems in the API documentation tools. The first alternative implementations for a minimal networked game on different modern high-level APIs studied here are written as a Javascript application and a combination of Actionscript (as3) for the client and Java for the server module. We have developed parsers for the internal / intermediate representations of class and method signatures in JsDoc JSON and AsDoc XML formats. The class information is read by a Python application to an internal model which contains the data for the OP calculation, implemented in another module in the same Python application.

To calculate the Message Points reflecting the *dynamic function calls*, we use the Closure Javascript compiler to traverse the source code to collect function calls and their

argument counts. A parser made with Python is used to read the function call data required to calculate the MPs.

While our tool calculates complete Class Point data, it currently omits two factors in Message Point data: the source and target counts of the interactions, and the novelty weight. While the tool tallies separate calls to each called function in the source code, it is not yet clear how to map them into the MP values. In Sobernig *et al.* the source and target counts were always set to 1. For the novelty weight we should check for each called operation called whether it is implemented in that class or inherited from a superclass. Our tool does not currently know the class of the object in which an operation is called. These omissions are not expected to affect the OP values significantly, at least not enough to affect the conclusions of this study.

Finally, to facilitate manual validation and visual communication of the data extracted from the source codes, our tool also creates UML class diagrams from the very same in-memory data structure that is used in the OP calculation. We chose the UXF format of the open source Umlet GUI diagram tool, due to its simple and straightforward XML format and the even simpler plaintext syntax used to describe individual UML elements, such as a class or a relation. This allows further manual editing of the diagrams with the GUI tool to improve the layout and annotation with notes.

Repository based automatic queries for OP analysis have been presented earlier by Henrich [12]. There a repository of documents, or abstract software design models (PCTE), were queried using the P-OQL language. We are not aware of any previous implementation of directly extracting data for the OP method from the source code.

IV. CASE STUDY IN API COMPLEXITY ANALYSIS

A. Pong as reference game

We propose using the Pong game as a minimal networked multiplayer reference game in the subsequent API complexity analysis. While Pong is tiny in its functionality, it is still sufficient for demonstrating key challenges in networked games, given the functional combination of the clients controlling their own paddles and the ball bouncing in the shared space. Pong has been used in networked game research earlier, recently in an interesting study of latency compensation techniques [13]. Also even a minimal game suffices to reveal the amount of software needed for the basic functionality: launching the networked software, establishing connections, handling players joining in and dropping out, and synchronizing gaming events.

B. Game platforms under comparison

In this section we briefly introduce the two game platforms whose API complexity we are going to compare, the open source realXtend Tundra SDK [4] and the Union, a proprietary closed source product¹. Both are relatively high-level platforms for networked games and bear several interesting similarities and differences for this study. Both are specifically designed for networking, which is exposed to the developer at

an abstract application level. That is, the developer and the game do not have know anything about sockets or network hosts. Instead, an abstract container object is provided (Room in Union, Scene in Tundra) and the game application logic listens to events from the container, for example when a new client joins the shared session/space. Also, both platforms provide an automated mechanism for synchronizing shared state over network. The shared state is stored in special attributes (objects of type Attribute) residing in the container (in Union directly in the Room object, in Tundra in the Components of the Entities in a Scene). The attributes are automatically shared among all participants, and notifications are provided for parties that have subscribed to be notified of changes. This way it is simple to for example set the game scores on the server, and show it in the client GUIs.

However, the two platforms also have fundamental differences and we discuss how they manifest in the implementation of the Pong game. TundraPong² was implemented by the leading author of this paper and an independent developer in two sessions totaling about six hours. UnionPong³ was downloaded from the Union website, where it is available as a tutorial example. While TundraPong is a script running atop the Tundra platform, UnionPong is a client application, to which networking functionality has been added by using Union's Reaktor Flash library. TundraPong utilizes a complete static scene datafile, where the game logic moves objects around. It runs on an existing client-server system, and utilizes several default components of the platform, most notably the data defining visual appearance and spatial instantiation and movement. In contrast, UnionPong not only has code to create the appearance of the game court (as it is called in Court.as), but also to define the data required for the spatial movement of an object (PongObject has x, y, direction, speed, width and height). In TundraPong, the predefined built-in Placeable component contains the position and the Rigidbody component contains shape information for collisions and speed vector for movement.

Thus, it is clear from the offset that UnionPong is more complex, due to the game code containing a much larger proportion of the implementation of the functionality. The upcoming API complexity analysis is still useful as it helps to answer the questions at hand: a) how the two APIs succeed in hiding the complexity from the developer and b) how our tool succeeds in evaluating the relative complexity of the two APIs.

C. Object-Point Data

The OP data of the two Pong implementations is presented in Table 1. As anticipated, TundraPong has clearly smaller OP values.

For TundraPong the OP data is extracted from the single Javascript source file (assets/game.js) that contains both client and server functionality in two respective classes, with GUI

TABLE I
OBJECT-POINT DATA FOR THE TWO PONG IMPLEMENTATIONS

metric	Tundra Pong		Union Pong		
	Full	Client only	Client Full	Client Net	Server
LoC	361	115	565	420	281
C	2	1	14	8	2
CP	75	27	180	140	75
MP	103	63	196	175	87
OP	178	90	376	315	162

and minimal game session management. UnionPong has separate client and server source code files in different languages using different libraries. Therefore, to facilitate more equal comparison, for TundraPong we also provide the OP data for the client only, even though it is included in the same source code file.

For UnionPong the OP data is calculated for all 14 client side ActionScript files and for selected 8 classes related to networking (GameManager, GameStates, KeyboardController, PongClient, PongObject, RoomAttributes, RoomMessages, UnionPong). The excluded classes cover GUI, the 2d scene implementation and general settings and utilities (clamp, ClientAttributes, Court, HUD, Rectangle and Settings). KeyboardController is included because it sends remote control messages from the player to the server (modifies client.paddle's attributes and says client.commit()).

UnionPong's Java server component (PongRoomModule.java) contains two classes: PongRoomModule (implements Module, Runnable) and PongObject, which is basically a duplicate of the same class in the client. As our OP data extraction tool does not yet support Java, we collected the OP data from the server component manually.

D. UML Diagrams

Figures 1 and 2 show the UML diagrams generated from the OP data by our tool for subsequent manual verification of the analysis and the API complexity.

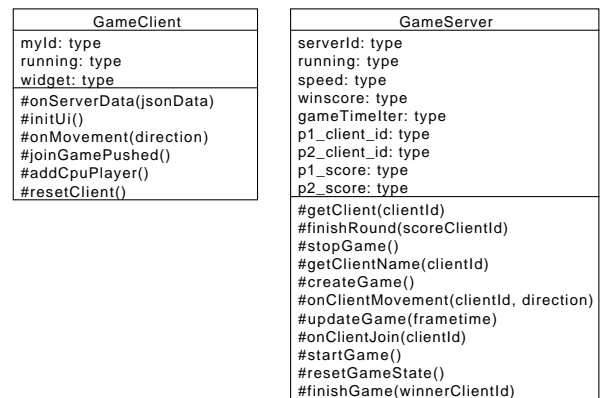


Fig. 1. The two classes in TundraPong game.js.

¹<http://www.unionplatform.com/>

²<https://github.com/realXtend/doc/tree/master/netgames/PongMultiplayer>

³http://www.unionplatform.com/?page_id=1229

HUD: flash.display.Sprite leftPlayerScore: type rightPlayerScore: type status: type #setLeftPlayerScore(score) #resetScores() #setRightPlayerScore(score) #setStatus(msg)	PongObject direction: type height: type speed: type width: type x: type y: type	GameStates GameManager joinTimer: type lastUpdate: type room: type state: type updateTimer: type #reset() #clientAttributeUpdateListener(e) #stopGameListener(fromClient) #removeRoomListeners() #setRoom(room) #joinTimerListener(e) #addRoomListeners() #roomAttributeUpdateListener(e) #updateBall(elapsed) #removePlayer(player) #addPlayer(player) #deserializeBall(value) #initPlayers() #roomJoinResultListener(e) #updatePlayer(player, elapsed) #roomJoinListener(e) #removeOccupantListener(e) #startGameListener(fromClient) #timerListener(e) #resetBall()
KeyboardController #setClient(client) #keyDownListener(e) #keyUpListener(e)	RoomMessages RoomAttributes ClientAttributes	
PongClient: reactor.CustomClient #getId() #init() #getIdPaddle() #updateAttributeListener(e) #commit() #deserializePaddle(value)	Settings Rectangle: flash.display.Sprite Court: flash.display.Sprite #showBall() #setBallPosition(x, y) #setRightPaddlePosition(x, y) #showLeftPaddle() #hideBall() #hideRightPaddle() #showRightPaddle() #hideLeftPaddle() #setLeftPaddlePosition(x, y)	
UnionPong: flash.display.Sprite reactor: type #initGame() #readyListener(e) #beginConnectListener(e)		

Fig. 2. The 13 classes in UnionPong client side ActionScript.

V. DISCUSSION

A. Interpretation and validity of Object-Point data

The fact that UnionPong obtains much higher OP values indicating a more complex API does not mean that the Union platform would be somehow inferior. Instead, it highlights the nature of game development at a different abstraction level. As we discussed in Section IV.B, the platforms achieve the basic functionality of the Pong game, such as synchronizing object movements and ball collisions and bounces, in different ways: UnionPong uses game specific messages, whereas TundraPong relies on the built-in functionality of the Tundra platform. Otherwise, the two APIs are very similar regarding networking. Both have an abstract container for the state, a Room in Union, and a Scene in Tundra. An application can store own custom state information as special attributes in the container, and the system takes care of automatically synchronizing changes to the state information. Both use callbacks heavily, for example to listen to new clients entering the service (an event of Room in Union's Reaktor and in the RoomModule on the Union server separately, an event of the Server core API object in Tundra server) and to attribute changes received from the network. They both also allow sending simple ad-hoc custom messages: Tundra uses them for game events such as informing of a victory with the associated data; UnionPong uses them for networking, including also paddle and ball movements, which Tundra does automatically. These similarities indicate that the OP analysis effectively captures the aforementioned differences in the scope and the abstraction level of the platforms.

Looking at the OP data and considering the OP method, our interpretation is that the OP method succeeds in illustrating the difference in scope and abstraction level between the two codebases. We have to ask whether the OP method does that better than some other, perhaps simpler, metrics would do. From previous research we know that the OP method does succeed in identifying complexity that a simple LoC metric would miss. Our data seems to support that same conclusion,

as the LoC measure would give a even larger complexity difference between the two implementations (115:565 for full clients). Based on qualitative analysis, we think that the smaller relative difference indicated by the OP data is more appropriate. Even though UnionPong client needs to do more, and especially has many more classes, most of the classes are very simple and most of the source code is not very complex. Considering only static class information, the difference would be even greater (27:180). TundraPong has relatively long methods and a lot of function calls, which lead to relative high MP count in the dynamic analysis (63:196). We think that changes the final OP score to a realistic ratio (90:376).

Based on the OP data we cannot really say whether the tool and the OP method misses something essential in the API complexity analysis. For example, the OP method does not take into account anything specific to networking: the need to think of connections, defining and sending network messages etc. They are of course accounted for as normal data definitions and function calls, but would some networking specific metric, for example for the number of messages, be more useful instead? Arguably, they present an additional complexity that the game developer has to manage.

B. Built-in platform logic vs. custom application logic

As TundraPong shows, implementing a game on an rich platform such as Tundra can require a comparatively small amount of work. However, custom game specific solutions for game object data, network messages, movement inter/extrapolation and collisions can easily be more powerful and even required. For example, if the game takes place on a small spherical world, a mini planet, Tundra's built-in euclidean movement techniques become suddenly much less useful. Therefore, the logic underlying the Union platform and other similar smaller libraries is sound. Game developers often need custom solutions, so the platform just provides the lower level tools for messaging and stays out of the way for the rest. There is a caveat, though. Optimizing and perfecting for example movement synchronization is not trivial, thus it is often useful to have a mature shared implementation. Tundra's rigidbody movement messaging was recently optimized, decreasing the message size from the 70 bytes/update of the initial naive design to 11 bytes/update in the current version. So having both reusable existing solutions and providing support for custom messaging makes sense. Tundra provides both, the high-level messaging similar to Union and used in the TundraPong game, and the low level custom messaging with the kNet library which is used for the built-in functionality as well.

C. Limitations of the study

The Pong game used as the minimal reference of a networked multiplayer game in this study is very simplistic. Much of the complexity of real networked games, and especially large scale commercial MMOGs, lies in areas not addressed by this study: service reliability, availability, restorability and scalability [1]. Networked programming in general is also

typically complex due to the need to handle several kinds of error situations, such as lost data, dropped connections and conflicts from simultaneous actions. The Pong reference implementations of this study may well be limited in that they do not handle such issues in the way a production quality game must handle, which probably increases the complexities. However, both Pong games are built on very high-level networked game platforms, which strive to take care of much of the complexity and hide it from the application layer. Whether and how they really achieve that cannot be determined from the data of this study, but would require a different analysis.

In future work the shortcoming of relying on too simple reference games in API complexity analysis could be addressed in two ways. First, we could analyze the codebases of real production quality games. However, typically a particular game exists only as a single implementation, which prevents comparative analysis. Nevertheless, the analysis could still provide valuable insight into assessing the evolution of the complexity of the game, and the correlation of the complexity with the monetary expenses of the development effort. Second, we could develop a more complex reference game and ensure its completeness. Such a reference game should be carefully specified to cover all relevant areas of networked gaming, but still remain small enough to allow implementation within a realistic timeframe. Existing canonical implementations may provide a starting point, as for example several commercial networked 3D first-person shooter (FPS) games have been open sourced (Quake, Cube2), and at least one highlevel platform already features a FPS as a tutorial (Torque3D).

VI. CONCLUSION

We first presented a tool for the automated quantitative assessment of the complexity of an API by extracting Sneed's Object-Point data from a source code using the API. We then compared the API complexity of two platforms for developing networking games, by extracting the OP data from two pre-existing implementations of the Pong game atop these platforms with the tool. The OP data and the related UML diagrams indicate that the tool successfully quantifies the different amounts of complexity that the game developer has to cope with on these two platforms. However, the difference in the relative complexity of the APIs is very much due to the different underlying design rationale of the platforms, the other relying on built-in functionality of the platform and the other on custom application logic embedded in the client.

In any case, automating data extraction for subsequent OP analysis opens up fascinating opportunities for future work in platform and API development. We could conduct longitudinal assessment of the evolution of the complexity of a particular API and codebase over time, dissect a software by running a series of OP analyses to pinpoint potential sources of complexity, and compare networking stacks based on different protocols for similar functionality.

VII. DOWNLOADS

Our tool and the datasets used in this study are available online at <https://github.com/realXtend/doc/tree/master/netgames/tools/>. The work-in-progress executable is pointcounter.py that contains the implementation of the Object-Point method.

ACKNOWLEDGEMENTS

To be added.

REFERENCES

- [1] T.-Y. Hsiao and S.-M. Yuan, "Practical middleware for massively multiplayer online games," *IEEE Internet Computing*, vol. 9, no. 5, pp. 47–54, 2005. [Online]. Available: <http://doi.ieeecomputersociety.org/10.1109/MIC.2005.106>
- [2] M. Henning, "Api design matters," *Commun. ACM*, vol. 52, no. 5, pp. 46–56, May 2009. [Online]. Available: <http://doi.acm.org/10.1145/1506409.1506424>
- [3] M. Cataldo and C. R. de Souza, "The impact of api complexity on failures: An empirical analysis of proprietary and open source software systems," [Online]. Available: <http://reports-archive.adm.cs.cmu.edu/anon/isr2011/CMU-ISR-11-106.pdf>
- [4] T. Alatalo, "An entity-component model for extensible virtual worlds," *IEEE Internet Computing*, vol. 15, no. 5, pp. 30–37, 2011.
- [5] S. Sobernig, P. Gaubatz, M. Strembeck, and U. Zdun, "Comparing complexity of api designs: an exploratory experiment on dsl-based framework integration," *SIGPLAN Not.*, vol. 47, no. 3, pp. 157–166, Oct. 2011. [Online]. Available: <http://doi.acm.org/10.1145/2189751.2047890>
- [6] M. Barth, "Api evaluation – an overview of api evaluation techniques," [Online]. Available: <http://dev.roleplaytalk.net/files/publications/api-evaluation.pdf>
- [7] J. Gerken, H.-C. Jetter, M. Zöllner, M. Mader, and H. Reiterer, "The concept maps method as a tool to evaluate the usability of apis," in *Proceedings of the 2011 annual conference on Human factors in computing systems*, ser. CHI '11. New York, NY, USA: ACM, 2011, pp. 3373–3382. [Online]. Available: <http://doi.acm.org/10.1145/1978942.1979445>
- [8] U. Farooq, L. Welicki, and D. Zirkler, "Api usability peer reviews: a method for evaluating the usability of application programming interfaces," in *Proceedings of the 28th international conference on Human factors in computing systems*, ser. CHI '10. New York, NY, USA: ACM, 2010, pp. 2327–2336. [Online]. Available: <http://doi.acm.org/10.1145/1753326.1753677>
- [9] C. R. B. de Souza and D. L. M. Bentolila, "Automatic evaluation of api usability using complexity metrics and visualizations," in *ICSE Companion*, 2009, pp. 299–302.
- [10] H. M. Sneed, "Estimating the costs of software maintenance tasks," in *ICSM*. IEEE Computer Society, 1995, pp. 168–181. [Online]. Available: <http://dblp.uni-trier.de/db/conf/icsm/icsm1995.html#Sneed95>
- [11] S. Thummalapenta and T. Xie, "Spotweb: Detecting framework hotspots and coldspots via mining open source code on the web," in *Proceedings of the 2008 23rd IEEE/ACM International Conference on Automated Software Engineering*, ser. ASE '08. Washington, DC, USA: IEEE Computer Society, 2008, pp. 327–336. [Online]. Available: <http://dx.doi.org/10.1109/ASE.2008.43>
- [12] A. Henrich, "Repository based software cost estimation," in *DEXA*, ser. Lecture Notes in Computer Science, A. Hameurlain and A. M. Tjoa, Eds., vol. 1308. Springer, 1997, pp. 653–662. [Online]. Available: <http://dblp.uni-trier.de/db/conf/dexa/dexa97.html#Henrich97>
- [13] F. E. Greger Wikstrand, Lennart Schedin, "High and low ping and the game of pong: Effects of delay and feedback," 2006. [Online]. Available: www.cs.umu.se/research/reports/2006/041/part1.pdf