

A comparison of networked game development APIs

This document was generated on 2012-08-30 at 20:13.

Networked game programming is complex compared to single player games. Several libraries and platforms exist to ease the development. In this exploratory study, we propose a technique for evaluating how much a networked game development platform succeeds in hiding complexity. We apply Sneed's Object-Points (OP) analysis to two pre-existing implementations of the same minimal networked multiplayer game: Pong. The OP technique successfully illustrates the different amounts of complexity the developer has to manage on the two alternative platforms. We have automated the source-code based analysis process, and suggest using it both for longitudinal studies of API development and for comparing alternative API approaches.

Introduction

Higher level abstractions in software are a common way to attempt to ease application development. Regarding networking, libraries exist to simplify managing connections and messaging. On a even higher level, distributed object systems automate remote calls and data synchronization. For an application developer, these systems are provided as a set of abstractions forming the application development interface (API).

It has been noted how making good APIs is hard -- and that creating a bad one is easy [api-matters](#). Even a small quirk in an API can accumulate to substantial problems in larger bodies of application code. API design has a significant impact on software quality, and increased API complexity is associated with increased software failure rate [cmu-api_failures](#).

An entity system for networked application development has been put forth in [Alatalo2011]. Developed in the open source Tundra SDK, it strives to apply best practices from game engine design literature, notably the aggregation using entity-component model. Specifically for networking, it features attribute autosynchronization, a simple form of transparent remote procedure calls (entity actions) and efficient customized movement messages with inter- and extrapolation logic (dead reckoning). The purpose of the abstract entity model, and the whole concrete platform, is to make multiplayer game development easy and productive. The goal in this study is to evaluate whether and how the Tundra API, and with it a few common practices in modern game and networking libraries, succeed in that.

How can a conceptual design of an entity system be really evaluated? How can we know how well a platform supports actual networked game development? These are not easy questions, but the answers would really help us concretely in game and platform development. We do not claim to provide final answers to all of it here. The area of software and API complexity analysis has however made interesting progress recently [api-complexity-analysis](#) [cmu-api_failures](#). By applying a software complexity analysis technique, we investigate one particular aspect of the quality of networked application platforms: the API complexity for a networked game developer.

We analyze API complexity by borrowing an approach from a previous study in a slightly different field. We conduct a comparative study of two alternative APIs for networked game development by analyzing the complexity of the same game implemented on the two platforms. The game is Pong, which is proposed as a minimal hello-world style example of a multiplayer game.

The article is organized as follows: Next, we provide background information on API complexity research, the selected game case and the alternative networked game development platforms. Then the conducting and the results of the Object-Points analysis is presented. Finally, results are discussed both to evaluate the applicability of the analysis method, and in light of explaining factors from the APIs.

Background

The research methodology - of API complexity research

Recently, software complexity analysis techniques have been applied to statistical (quantitative) studies of API complexity. [cmu-api_failures](#) studies 2 large corporate software projects and 9 open source projects and finds a link between API complexity and increases in failure proneness of the software (bug reports from the field). The masses of source code are quantified with measures such as API size and dispersion. Building on existing work, API complexity is calculated simply from the number of public methods and attributes. In the discussion it is noted how this is severely limited: for example, it fails to take into account pre- and post-invocation assumptions of the API and possibly required sequences of invocation [cmu-api_failures](#).

We do not have external statistics data from hours used for development or reported software failures of games to study. Also simplistic measures, suitable for analyzing large bodies of source code, would miss the subtle issues which arise in networked programming on a framework which attempts to hide the intricacies of networking from the application developer. It would be interesting to organize an experiment where a number of test teams develop the same networked game on alternative platforms, from the same specifications, and the development time and number of mistakes would be analyzed. That is however out of the scope here.

In a different approach, a study of 4 alternative implementations, on different frameworks, of the same application uses Object-Points (OP) analysis to quantify the code bases for the comparison [api-complexity-analysis](#). OP has originally been developed for estimating development effort, but there the authors adopt it to calculate the complexity of existing software for complexity comparisons. Number of classes, their members and operation calls are counted and assigned adjustment weights in the calculation. Intermediate UML models are used as the data source which allows comparing programs in different languages [api-complexity-analysis](#). This kind of fine grained OP analysis is applicable for our purposes here. It does not capture all the elements of API complexity, but gives useful metrics for comparisons.

The game of Pong

We propose using Pong as a minimal networked multiplayer game. It is tiny in functionality, but still demonstrates key issues with networking and games with the combination of the clients controlling their own paddles and the ball bouncing in the shared space. Pong has been used in networked game research earlier, recently in an interesting study of latency compensation techniques [pong-ping](#). Also even a minimal game suffices to reveal the amount of software needed for all the basics: establishing connections, handling players joining in and dropping out, and just getting the networked software up and running.

For further studies, devising a set of different kind of small games, and perhaps some larger sufficiently complex game, would really allow rich comparative API analysis.

Platforms: realXtend Tundra SDK and Union Platform

For this initial study, we selected two relatively high-level networked game platforms: realXtend Tundra SDK (open source) and the Union Platform (closed source proprietary). They bear several key similarities and differences which are interesting for the study:

Both Tundra and Union are specifically for networking, and expose it to the developer on an abstract application level. That is, the games do not know anything about sockets or network hosts. Instead, an abstract container object is provided (Room in Union, Scene in Tundra). Application logic listens to events from the container, for example when a new client joins the shared session/space.

Also, both platforms provide an automated mechanism for synchronizing state over the network. The shared state is in special attributes (objects of type Attribute), which are in the container (in Union directly in the Room object, in Tundra in entities in the Scene). The attributes are automatically shared among all the participants, and provide events for interested parties to get notified of changes. This way it is simple

to for example set the game score points on the server, and show it in the GUI in clients.

However, there is one fundamental difference in the platforms and how they are used in the Pong examples studied here. TundraPong is a script running on the Tundra platform. UnionPong is a new client application, to which the networking has been added by using Union's Reaktor Flash library. The Tundra game utilizes a complete static scene datafile where the game logic just starts moving objects around. It runs on an existing client-server system, and utilizes several default components from the platform: notably all the data for the appearance and spatial instancing. In contrast, UnionPong not only has code to create the appearance of the game court (as it is called in Court.as), but also to define what data is required for a spatial moving object (PongObject has x, y, direction, speed, width and height). Tundra, again, has the position in the builtin predefined Placeable component and the size and shape information for collisions, and the speed vector for movement, in the physics module's Rigidbody component. Also with networking there is a great difference: OnionPong sends own custom movement messages for all the movement, and has also custom server side code to do ball bouncing, whereas on Tundra the default movement replication and physics collisions are used.

So it is clear at the start that UnionPong is more complex, due to having much more of the implementation in the game/application code. The analysis is still interesting as it helps to answer the questions at hand: a) how much do the alternative APIs manage to hide complexity and b) how well does the selected analysis technique apply to networked game API evaluation.

For more results, at least these two additional Pong implementations should be added to the analysis in future work:

1. An alternative TundraPong style game where the defaults from an underlying platform are used to the fullest, for example with the Unreal engine.
2. A version made with a different networked programming paradigm, such as the Emerson language which is a Javascript variant by the Sirikata project for networked applications, without attribute autosynchronization but using messaging exclusively instead [sirikata-scripting](#).

The analysis here is limited to the two platforms simply because we do not have more implementations (Pong source codes) to study yet. The Tundra one was initiated by the author (only the scene and trivial computer opponent logic as a test), and later completed by an independent developer (he made all the networking and game control code). The Union one we found with an Internet search.

Application of Object-Point analysis

The chosen Sneed's Object-Point (OP) analysis was conducted by automating the collection of most of the key data to derive the variables in the equation. We apply the technique following what has been used for API complexity analysis before in [api-complexity-analysis](#)

To read the *static class data* for the **Class Points** (CP), we utilize existing source code parsing and annotation systems in API documentation tools. The first alternative implementations of a minimal networked game on different modern high-level APIs studied here are written as a a) Javascript application and b) a combination of Actionscript (as3) for the client and Java for the server module. We developed parsers for the internal / intermediate representation of class and method signatures of JsDoc JSON and AsDoc XML. (The single Java class for b) server we may analyze manually). The class information is read in a Python application to an internal model which contains the data for the Sneed points calculation, implemented in another module in the same Python application.

For the *dynamic function call* information, to calculate the **Message Points** (MP) in the overall OP analysis, we use the Closure Javascript compiler to traverse the source code to collect function calls and their argument counts. Basic filtering with AWK is used to filter in the relevant information from the Closure tree. To be able to analyze also Actionscript code, we do text processing to strip AS extensions to the basic ECMA/Javascript (remove public/private definitions and type declarations). A simple parser made with Python is used to read the function call data required to calculate MPs. This completes the automated data collection and processing developed for the OP calculations here.

The software to run the calculations, together with the datasets used in the analysis here, is available from <https://github.com/realXtend/doc/tree/master/netgames/tools/> (pointcounter.py is the executable, with the

formula for OP = CP + MP).

TODO: add the equation + legend here

Results

Full applications

Tundra PongMultiplayer

game.js: 74 OP (CP + MP)

UnionPlatform Pong tutorial

client 14x .as3: 220.5 OP (CP + MP)

- UnionPong/Java/PongRoomModule.java

Only the networking code

- Selected classes, explain the criteria.

unionpongnet = ["GameManager", "GameStates", "KeyboardController", "PongClient", "PongObject", "RoomAttributes", "RoomMessages", "UnionPong"]

KeyboardController is included because it is exactly what sends the remote control messages from the player to the server (modifies client.paddle's attributes and says client.commit()).

client 8x .as: 147.0

Discussion

How should we interpret this result? There are several things to consider, these are visited in the following: 1. validity of the analysis technique, the automated (partial) Object-Point analysis 2. nature, suitability and use of scripting vs. application development libraries 3. observations of the high-level network programming APIs studied here. 4. limitations: the many areas of analysis outside the focus here (scalability, efficiency of the networking etc)

1. Validity of the analysis

We apply Sneed's Object-Point analysis, following how it has been adopted to API complexity evaluation in [api-complexity-analysis](#), as closely as we could with the automated source code analysis. The validity must thus be evaluated from two viewpoints: a) applicability of OPs to API complexity analysis in general and b) the deviations from the intended calculation due to limits of the analysis software.

The OP sums of the full examples have an order of magnitude (right? XXX) sized difference in the proposed complexity of the two implementations of the same game. Noting the aforementioned substantial difference in the nature and scope of the implementations, the ratio of 74:273 (XXX fix when nums update) seems correct for codebases of 2 sizeable and 14(+1) mostly small classes respectively.

TODO: what was left out from analysis (was anything, in the end? XXX)

2. On scripting vs own client development

TODO

- as the data points out, implementing something on an existing platform can be comparatively very little work

- making an own application (client) is easily powerful and straightforward for own custom things, however
- same existing modules/components can be used either way, though. still simpler when don't need to deal with application init and connecting etc.
- does the complexity lurk somewhere still?

3. Observations of the high-level network programming APIs

The APIs under study here are very similar regarding the networking. They both have an abstract container for the state: a Room in Union, and a Scene in Tundra. Application can put own custom state information as special attributes in that container, and the system takes care of automatically synchronizing changes to that data.

Both use callbacks heavily, for example both to listen to new clients entering the service (an event of Room in Union's Reaktor and in the RoomModule on the Union server separately, an event of the Server core API object in Tundra on server side) and to attribute changes coming in over the network.

They both also allow sending simple ad-hoc custom messages, which the Tundra version uses for game events such as informing of a victory (with the associated data), and UnionPong uses for all networking (also paddle and ball movements).

With this in mind, we would expect the difference in the complexity sum derive from the scope of the implementations used in the analysis.

TODO: return to this when the numbers from network-code-only analysis are in too?!?

4. Limitations

the many areas of analysis outside the focus here (scalability, efficiency of the networking etc)

TODO

Conclusions

TODO

(We are happy and curious about using this tool for many kinds of comparisons: longitudinal studies of a single API over time, comparisons of e.g. networking stacks when using different protocols for similar functionality, ... or?)

Similarities and differences of using a platform as ready made client software, on which just run scripts, vs. libraries to create own applications, are interesting to study more. Same software components (libraries, modules etc) can be used in both configurations -- what is more suitable may well depend on the particular case.

(XXX Q: where does complexity lurk? should we consider the leaks here? does Onion have something to handle them? at least had the Attribute setting exception in the java server XXX)

References

cmu-api_failures(1, Marcelo Cataldo¹, Cleidson R.B. de Souza² (2011). The Impact of API Complexity on Failures: An Empirical Analysis of Proprietary and Open Source Software Systems. <http://reports-archive.adm.cs.cmu.edu/anon/isr2011/CMU-ISR-11-106.pdf>
 2, 3, 4)

api-complexity-analysis Comparing Complexity of API Designs: An Exploratory Experiment on DSL-based Framework Integration.
<http://www.sba-research.org/wp-content/uploads/publications/gpce11.pdf>

pong-ping High and Low Ping and the Game of Pong. <http://www.cs.umu.se/~greger/pong.pdf>

sirikata-scripting Bhupesh Chandra, Ewen Cheslack-Postava, Behram F. T. Mistree, Philip Levis, and David Gay. "Emerson: Scripting for Federated Virtual Worlds", Proceedings of the 15th International Conference on Computer Games: AI, Animation, Mobile, Interactive Multimedia, Educational & Serious Games (CGAMES 2010 USA).
<http://sing.stanford.edu/pubs/cgames10.pdf>