API complexity analysis of networked game platforms using a reference game: case of Pong

Abstract—Networked game programming is complex compared to single player games. Several libraries and platforms exist to ease the development. In this exploratory study, we propose a technique for evaluating how much a networked game development platform succeeds in hiding complexity. We apply Sneed's Object-Points (OP) analysis to two pre-existing implementations of the same minimal networked multiplayer game: Pong. The OP technique succesfully illustrates the different amounts of complexity the developer has to manage on the two alternative platforms. We have automated the source-code based analysis process, and suggest using it both for longitudal studies of API development and for comparing alternative API approaches.

I. INTRODUCTION

Networked application programming is generally more complex than standalone software development. The developer typically needs to deal with events, conflicts and error conditions originating from other parts of the distributed system as well as the local user. This is emphasized in multiuser real-time systems compared to the relatively leisurely request-response interaction patterns of most client-server applications. Developing a massively multiplayer online game (MMOG) has been stated to typically take two to three times as long as creating and launching a single-player game [middleware].

Higher level abstractions in software are a common way to attempt to ease application development. Regarding networking, libraries exist to simplify managing connections and messaging. On a even higher level, distributed object systems automate remote calls and data synchronization. For an application developer, these systems are provided as a set of abstractions forming the application development interface (API).

It has been noted how making good APIs is hard -- and that creating a bad one is easy [api-matters]. Even a small quirk in an API can accumulate to substantial problems in larger bodies of application code. API design has a significant impact on software quality, and increased API complexity is associated with increased software failure rate [cmu-api_failures].

An entity system for networked application development has been put forth in [Alatalo2011]. Developed in the open source Tundra SDK, it strives to apply best practices from game engine design literature, notably the aggregation using entity-component model. Specifically for networking, it features attribute autosynchronization, a simple form of transparent remote procedure calls (entity actions) and efficient customized movement messages with inter- and extrapolation logic (dead reckoning). The purpose of the abstract entity model, and the whole concrete platform, is to make multiplayer game development easy and productive. The goal in this study is to evaluate whether and how the Tundra API, and with it a a few common practices in modern game and networking libraries, succeed in that.

How can a conceptual design of an entity system be really evaluated? How can we know how well a platform supports actual networked game development? These are not easy questions, but the answers would really help us concretely in game and platform development. One presentation of a MMOG middleware proposes four essential "ease of" requirements: ease of development, deployment, maintenance and change. However in the evaluation they note the difficulty of quantitative measurement of e.g. ease of development or change, and review platform scalability only [middleware]. Here we focus on the difficult ease-of question instead. We do not claim to provide final answers to all of it here. The area of software and API complexity analysis has however made interesting progress recently [api-complexity-analysis], [cmu-api_failures]. By applying a software complexity analysis technique, we investigate one particular aspect of the quality of networked application platforms: the API complexity for a networked game developer.

We analyze API complexity by following an approach from a previous study in a slightly different field [api-complexity-analysis]. We conduct a comparative study of two alternative APIs for networked game development by analyzing the complexity of the same game implemented on the two platforms. The game is Pong, which is proposed as a minimal hello-world style example of a multiplayer game.

The article is organized as follows: Next, we provide background information on API complexity research, the selected game case and the alternative networked game development platforms. Then the conducting and the results of the Object-Points analysis is presented. Finally, results are discussed both to evaluate the applicability of the analysis method, and in light of explaining factors from the APIs.

II. BACKGROUND

A. The research methodology - of API complexity research

Recently, software complexity analysis techniques have been applied to statistical (quantitative) studies of API complexity. [cmu-api_failures] studies 2 large corporate software projects and 9 open source projects and finds a link between API complexity and increases in failure proneness of the software (bug reports from the field). The masses of source code are quantified with measures such as API size and dispersion. Building on existing work, API complexity is

calculated simply from the number of public methods and attributes. In the discussion it is noted how this is severely limited: for example, it fails to take into account pre- and post-invocation assumptions of the API and possibly required sequences of invocation [cmu-api_failures].

More generally, API usability research has recently gotten attention in human centric research. Both traditional HCI usability evaluation techniques have been adapted to API evaluation, and also novel approaches specific to API usability research have been put forth [conceptmaps]. Traditional techniques include the think aloud protocol, heuristic evaluation and cognitive walkthoughs [overview]. Metrics in those studies include for example the completion times of predefined tasks. Human observation based API evaluation is challenging due to several reasons, compared to usability studies of simpler tasks accomplished with GUIs. Developing even a small application takes easily weeks so it is difficult to fit valid tasks in a typical 1-2 hour observation session [conceptmaps].

Novel approaches developed especially for API evaluation and development include peer reviews [apipeerreview], walkthroughs and the concept maps method [conceptmaps], [overview]. These avoid the problems of traditional HCI methods, notably by involving real world usage of the API over a long period of time. They are still considerably laborsome and essentially qualitative analysis. We plan to conduct such studies in future work on network game platform development.

Here, however, we investigate the suitability of quantitative and even automated API complexity measures. We base the analysis on existing bodies of source code, which has several advantages: real world data (source codes of existing applications) can be applied, and the analysis can be quick being not very laborsome and giving immediate feedback. Longitudinal studies of API development over time can be straightforward to conduct by running the analysis for different versions of the software. If fully automated, the analysis can be even run as a part of a continuous integration setup.

To address the issue of too simplistic software metrics (such as lines of code count), we apply the relatively sophisticated Object-Points (OP) analysis by Sneed. It has been proposed for API complexity evaluation in [api-complexity-analysis], where four alternative implementations, on different frameworks, are used for a comparative analysis. OP has originally been developed for estimating development effort, but there the authors adopt it to calculate the complexity of existing software for complexity comparisons. Number of classes, their members and the set of operations called are counted and assigned adjustment weights in the calculation. A key decision is to apply a surrogate measurement: the APIs are not analyzed directly, but programs developed against them are analyzed. The focus is on how a API is used, which is in line with the practice in the aforementioned HCI studies. Measuring the complexity, for example the sizes, of the APIs themselves would give misleading results, as a more complete API would appear more complex, even if it provided good concepts so that the task at hand could be in fact accomplished with just a small subset of the full API. In OP analysis, intermediate UML models are used as the data source which allows comparing programs in different languages [api-complexity-analysis]. Importantly, the Sneed measure allows direct tracking from indicator values to program structures. This is elemental for the purposes of API evaluation and design -- for example if many codebases get a high proportion of their complexity value due to a specific part of the API, it can then be examined qualitatively.

B. The game of Pong

We propose using Pong as a minimal networked multiplayer game. It is tiny in functionality, but still demonstrates key issues with networking and games with the combination of the clients controlling their own paddles and the ball bouncing in the shared space. Pong has been used in networked game research earlier, recently in an interesting study of latency compensation techniques [pong-ping]. Also even a minimal game suffices to reveal the amount of software needed for all the basics: establishing connections, handling players joining in and dropping out, and just getting the networked software up and running.

For further studies, devising a set of different kind of small games, and perhaps some larger sufficiently complex game, would really allow rich comparative API analysis.

C. Platforms: realXtend Tundra SDK and Union Platform

For this initial study, we selected two relatively high-level networked game platforms: realXtend Tundra SDK (open source) and the Union Platform (closed source proprietary). They bear several key similarities and differences which are interesting for the study:

Both Tundra and Union are specifically for networking, and expose it to the developer on an abstract application level. That is, the games do not know anything about sockets or network hosts. Instead, an abstract container object is provided (Room in Union, Scene in Tundra). Application logic listens to events from the container, for example when a new client joins the shared session/space.

Also, both platforms provide an automated mechanism for synchronizing state over the network. The shared state is in special attributes (objects of type Attribute), which are in the container (in Union directly in the Room object, in Tundra in Components of the Entities in a Scene). The attributes are automatically shared among all the participants, and provide events for interested parties to get notified of changes. This way it is simple to for example set the game score points on the server, and show it in the GUI in clients.

However, there is one fundamental difference in the platforms and how they are used in the Pong examples studied here. TundraPong is a script running on the Tundra platform. UnionPong is a new client application, to which the networking has been added by using Union's Reaktor Flash library. The Tundra game utilizes a complete static scene datafile where the game logic just starts moving objects around. It runs on an existing client-server system, and utilizes several default components from the platform: notably all the data for the appearance and spatial instancing. In contrast, UnionPong not

only has code to create the appearance of the game court (as it is called in Court.as), but also to define what data is required for a spatial moving object (PongObject has x, y, direction, speed, width and height). Tundra, again, has the position in the builtin predefined Placeable component and the size and shape information for collisions, and the speed vector for movement, in the physics module's Rigidbody component. Also with networking there is a great difference: OnionPong sends own custom movement messages for all the movement, and has also custom server side code to do ball bouncing, whereas on Tundra the default movement replication and physics collisions are used.

So it is clear at the start that UnionPong is more complex, due to having much more of the implementation in the game/application code. The analysis is still interesting as it helps to answer the questions at hand: a) how much do the alternative APIs manage to hide complexity and b) how well does the selected analysis technique apply to networked game API evaluation.

For more results, at least these two additional Pong implementations should be added to the analysis in future work:

1. An alternative TundraPong style game where the defaults from an underlying platform are used to the fullest, for example with the Unreal engine.

2. A version made with a different networked programming paradigm, such as the Emerson language which is a Javascript variant by the Sirikata project for networked applications, without attribute autosynchronization but using messaging exclusively instead [sirikata-scripting].

The analysis here is limited to the two platforms simply because we do not have more implementations (Pong source codes) to study yet. Also we find that a careful review is in place first to evaluate the suitability of this kind of Object-Points analysis, before continuing to apply it more. The Tundra one was initiated by the author (only the scene and trivial computer opponent logic as a test), and later completed by an independent developer (he made all the networking and game control code). The Union one we found with an Internet search.

III. APPLICATION OF OBJECT-POINT ANALYSIS

The chosen Sneed's Object-Point (OP) analysis was conducted by automating the collection of most of the key data to derive the variables in the equation. We apply the technique following what has been used for API complexity analysis before in [api-complexity-analysis]. Here we give a brief overview of Sneed's OP analysis itself, and describe how we derive the data from source code analysis.

A. Sneed's Object-Point analysis

(NOTE: this is a little new background treatment again - consider moving some of this to 2. etc XXX)

Software cost estimation has been of paramount importance in the field of software engineering, and various approaches have been developed for it through the decades. The early COCOMO model uses simply program size (lines of code) to estimate development effort, but later the Function-Point, Data-Point and finally Object-Point methods base the analysis on functionality and other properties of the program [henrich97repositorybased]. Recently the Object-Point (OP) method has been used for analysing existing implementations, for API complexity comparison purposes, even though it was originally developed for early work estimate analysis based on UML design diagrams [api-complexity-analysis]. Arguably, it is rich enough to explore structural and dynamic properties of software for meaningful complexity data.

For example in the preceeding API complexity analysis OP study that we follow here, two of the four compared implementations would get the opposite results in a simplistic lines of code (LOC) analysis. That is, the PHP implementation there features only 48 LOC but results in 356.34 OP, whereas the domain specific language (DSL) version is 144 LOC and 266.76 OP [api-complexity-analysis]. Their explanation is that "an API user is only exposed to an API feature chunk of low structural complexity", as the chunk's "size is limited in terms of participating classes and the smallest number of operations per class" and it "shows a relatively weak connectedness of classes (H = 1), resulting from the small number of associations and generalizations between the classes".

That is of utmost importance to our interest in making networked game development easier with a good API. We are after a limited set of good concepts with clear interactions that a game developer could learn easily and grow to master. Not all lines of code are equal -- a bad API makes it a struggle to get even a few operations working if the developer has to hunt for functionality that is scattered around in an incoherent way.

The Object-Points, as applied here, are a sum of two parts: Class Points (CP) and Message Points (MP).

Class points, CP is calculated from the static class structure, specifically: the class count and sums of attribute, operation and relation counts. Weights are used to correct the values for the overall calculation. Class inheritance is taken into account by calculating novelty weights for specializing classes.

Message points, MP is defined by the set of operations (functions/methods) *actually used* in the software. First, the number of operations is used. Then the parameter count for each called operation is collected. Also the source and target counts of the operation calls are established. Again, novelty weights are used to compensate for repeated occurrences due to subclassing.

TODO: add the equation + legend here -- but refer to the other paper for more, or do we need to explain every detail here too?

B. Reading class and interaction data from source code

To read the *static class data* for the **Class Points** (CP), we utilize existing source code parsing and annotation systems in API documentation tools. The first alternative implementations of a minimal networked game on different modern high-level

APIs studied here are written as a a) Javascript application and b) a combination of Actionscript (as3) for the client and Java for the server module. We developed parsers for the internal / intermediate representation of class and method signatures of JsDoc JSON and AsDoc XML. (XXX The single Java class for b) server we may analyze manually). The class information is read in a Python application to an internal model which contains the data for the Sneed points calculation, implemented in another module in the same Python application.

For the *dynamic function call* information, to calculate the **Message Points** (MP) in the overall OP analysis, we use the Closure Javascript compiler to traverse the source code to collect function calls and their argument counts. To be able to analyze also Actionscript code, we do text processing to strip AS extensions to the basic ECMA/Javascript (remove public/private definitions and type declarations). A parser made with Python is used to read the function call data required to calculate MPs. This completes the automated data collection and processing developed for the OP calculations here.

Finally, to facilitate manual validation and visual communication of the data mined from the source codes, we added functionality to create UML class diagrams from the very same in-memory data structure which is for the OP calculation. We chose the UXF format of the open source Umlet GUI diagram tool, due to it's simple and straightforward XML document format and the even simpler plaintext syntax used to describe the individual UML elements, such as a class or a relation. It is useful to be able to manually edit the diagrams further with the GUI tool to improve the layout and add notes.

All this software to run the calculations, together with the datasets used in the analysis here, is available from https://github.com/realXtend/doc/tree/master/netgames/tools/ (point-counter.py is the executable, with the implementation of the equation).

Repository based automatic queries for OP analysis have been presented earlier in [henrich97repositorybased]. There a repository of *documents*, or abstract software design models (PCTE) is queried for automatic OP calculations using the P-OQL language. We are not aware of previous implementations of deriving data for OP calculations from source code only. Automating the calculation opens up fascinating possibilities for platform and API development in future work, such as longitudal evaluation of API complexity evolution, and dissecting a body of software by running a series of calculations to pinpoint potential sources of complexity.

IV. RESULTS

The results for the Object Points analysis for the two codebases are presented in table 1.

For TundraPong, the single Javascript source file (assets/game.js) is included. It features both client and server functionality in two classes respectively. It is the complete implementation with GUI and the minimal game session management. The Union implementation has the client and server separately, in different languages and made with different

TABLE I
OBJECT-POINTS FOR TUNDRA AND UNION PONGS.

metric	Tundra Pong full	Tundra Pong client only	Union Pong Client full	Union Pong Client net only	Union Pong Server
Lines of code	361	115	565	420	281
Number of classes	2	1	14	8	2
Class Points	75	27	180	140	-
Message Points	103	63	196	175	-
Object Points	178	143	376	315	-

libraries. Therefore to enable more equal comparisons we include the metrics for the client side code only also for the Tundra implementation, even though it is included in the same source code file.

For UnionPong, all the client side ActionScript files (14) are included for the full run, and selected 8 for the network code only calculation. The selection is made on the class level: the classes which involve networking are included in full, not edited line-by-line to include networking code only. The included classes are: GameManager, GameStates, KeyboardController, PongClient, PongObject, RoomAttributes, RoomMessages, UnionPong. The excluded classes cover GUI, the 2d scene implementation and general settings and utilities, and are called: clamp, ClientAttributes, Court, HUD, Rectable and Settings. KeyboardController is included because it is exactly what sends the remote control messages from the player to the server (modifies client.paddle's attributes and says client.commit()).

The Java-written server side component for UnionPong, in PongRoomModule.java, features two classes: PongRoomModule (implements Module, Runnable) and PongObject, which is basically a duplicate of the same class on the client side. We have not conducted the OP analysis for it, and the automated system does not support Java currently. (XXX NOTE: data/unionpong/server.txt has the static data for CP analysis now, grepped from the source file, so that would be quite quick to do .. and MP could be done manually).

A. UML Diagrams

The data used in the calculations is also generated to UML class diagrams by the analysis software, for manual verification of the source code analysis process, and for (XXX thinking about the codebases & complexity calcs).

GameClient
myId: type
running: type
widget: type
#onServerData(jsonData)
#initUi()
#onMovement(direction)
#joinGamePushed()
#addCpuPlayer()
#resetClient()

GameServer
serverId: type
running: type
speed: type
winscore: type
gameTimeIter: type
p1_client_id: type
p2_client_id: type
p1_score: type
p2_score: type
#getClient(clientId)
#finishRound(scoreClientId)
#stopGame()
#getClientName(clientId)
#createGame()
#onClientMovement(clientId, direction)
#updateGame(frametime)
#onClientJoin(clientId)
#startGame()
#resetGameState()
#finishGame(winnerClientId)

Fig. 1. The two classes in TundraPong game.js.

HUD: flash.display:Sprite
leftPlayerScore: type
rightPlayerScore: type
status: type
#setLeftPlayerScore(score)
#resetScores()
#setRightPlayerScore(score)
#setStatus(msg)

KeyboardController
#setClient(client)
#keyDownListener(e)
#keyUpListener(e)

PongClient: reactor:CustomClient
#getSide()
#init()
#getPaddle()
#updateAttributeListener(e)
#commit()
#deserializePaddle(value)

UnionPong: flash.display:Sprite
reactor: type
#initGame()
#readyListener(e)
#beginConnectListener(e)

PongObject
direction: type
height: type
speed: type
width: type
x: type
y: type

RoomMessages

RoomAttributes

ClientAttributes

Settings

Rectangle: flash.display:Sprite

Court: flash.display:Sprite
#showBall()
#setBallPosition(x, y)
#setRightPaddlePosition(x, y)
#showLeftPaddle()
#hideBall()
#hideRightPaddle()
#showRightPaddle()
#hideLeftPaddle()
#setLeftPaddlePosition(x, y)

GameStates

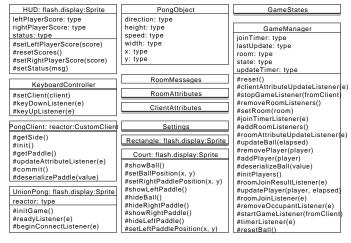| GameManager |
|---|
| joinTimer: type |
| lastUpdate: type |
| room: type |
| state: type |
| updateTimer: type |
| #reset() |
| #clientAttributeUpdateListener(e) |
| #stopGameListener(fromClient) |
| #removeRoomListeners() |
| #setRoom(room) |
| #joinTimerListener(e) |
| #addRoomListeners() |
| #roomAttributeUpdateListener(e) |
| #updateBall(elapsed) |
| #removePlayer(player) |
| #addPlayer(player) |
| #deserializeBall(value) |
| #initPlayers() |
| #roomJoinResultListener(e) |
| #updatePlayer(player, elapsed) |
| #roomJoinListener(e) |
| #removeOccupantListener(e) |
| #startGameListener(fromClient) |
| #timerListener(e) |
| #resetBall() |

Fig. 2. The 13 classes in UnionPong client side ActionScript.

V. Discussion

How should we interpret this result? There are several things to consider, these are visited in the following: A. validity of the analysis technique, the automated (partial) Object-Point analysis B. nature, suitability and use of scripting vs. application development libraries C. observations of the high-level network programming APIs studied here. D. limitations: the many areas of analysis outside the focus here.

A. Validity of the analysis

We apply Sneed's Object-Point analysis, following how it has been adopted to API complexity evaluation in [api-complexity-analysis], as closely as we could with the automated source code analysis. The validity must thus be evaluated from two viewpoints: a) applicability of OPs to API complexity analysis in general and b) the deviations from the intended calculation due to limits of the analysis software.

(XXX FIX and update: The OP sums of the full examples have an order of magnitude (right? XXX) sized difference in the proposed complexity of the two implementations of the

same game. Noting the aforementioned substantial difference in the nature and scope of the implementations, the ratio of 74:273 (XXX fix when nums update) seems correct for codebases of 2 sizeable and 14(+1) mostly small classes respectively.)

The fact that UnionPong gets much higher complexity points does not mean that the Union platform would be somehow bad. Instead, it highlights the nature of developing on a different abstraction level. As described already in the background information (section 2C), the implementations achieve the basic functionality --- such as synchronizing object movements and ball collisions and bounces --- in alternative ways: UnionPong with game specific messages, TundraPong automatically by relying on the built-in functionality of the underlying platform.

Looking at the data and considering the OP analysis technique, our understanding is that the analysis succeeds in illustrating this difference in the codebases. A question is whether OP analysis does that better than some other, perhaps simpler, metric would do. From previous research we know that OP does succeed in analyzing complexity that a simple lines of code (LoC) measure would miss (the case of PHP and DSL implementations in section 3A here, reviewing [api-complexity-analysis]). Our data seems to support that, as the LoC measure would give a much smaller complexity difference for the two implementations. Based on qualitative analysis, we think that the higher difference is correct (UnionPong needs to do much more, has many more classes and function calls). But based on the data here, we can not say whether utilizing OP was better than for example the simpler code size analysis used in [cmu-api_failures]. We think that the analysis of dynamic code behaviour, collecting the information about operations actually called for the Message-Points calculation, is essential. We cannot however support that reasoning with the data here, as the result would be similar with only the static Class-Points as well. This may be due to the similar style of the codebases. We rely on previous research for the generic validation of the OP construct.

Also we can not tell from this data whether something essential was missed in the analysis. For example, the technique does not take into account anything specific to networking: the need to think of connections, defining and sending network messages etc. They are of course accounted for as normal data definitions and function calls, but would some specific measure for example for the number of network messages be useful instead? Arguably they are additional complexity that the developer has to manage.

TODO: what was left out from analysis (was anything, in the end? XXX)

B. On scripting vs. own client development

As the data points out, implementing something on an existing platform can be comparatively very little work. (XXX NOTE: own game specific custom things can be powerful -- however, for common tasks, a generic optimized solution in a platform can work better (the optimized movement messages

in Tundra -- a change to use the measurement data as well? !!! XXX)

- making an own application (client) is easily powerful and straightforward for own custom things, however
- same existing modules/components can be used either way, though. still simpler when don't need to deal with application init and connecting etc.
- does the complexity lurk somewhere still?

C. Observations of the high-level network programming APIs

The APIs under study here are very similar regarding the networking. They both have an abstract container for the state: a Room in Union, and a Scene in Tundra. Application can put own custom state information as special attributes in that container, and the system takes care of automatically synchronizing changes to that data.

Both use callbacks heavily, for example both to listen to new clients entering the service (an event of Room in Union's Reaktor and in the RoomModule on the Union server separately, an event of the Server core API object in Tundra on server side) and to attribute changes coming in over the network. (NOTE: both have the attr-changed granularity prob! XXX)

They both also allow sending simple ad-hoc custom messages, which the Tundra version uses for game events such as informing of a victory (with the associated data), and UnionPong uses for all networking (XXX these were said already: (also paddle and ball movements, which Tundra does automatically).

With this in mind, we would expect the difference in the complexity sum derive from the scope of the implementations used in the analysis.)

D. Limitations

The Pong game examples studied here are minimal tutorial like simplistic implementations. Much of the complexity of real networked games, and especially large scale commercial MMOs, lies in areas this example does not address: service reliablity, availability, restorability and scalability are listed as vital issues in [middleware]. Networked programming in general is also typically complex due to the need to handle several kinds of error situations, such as lost data, dropped connections and conflicts from simultaneous actions.

The current reference implementations here may well not be complete in the sense that they would handle such issues in the way a production quality game must, which probably would increase the complexities. However, they both are built on very high-level networked game platforms, which strive to take care of much of the complexity and hide it from the application layer. Whether and how they really achieve that is not known based on the data here, but would require a different analysis.

The basic issue of too simple example games can be addressed in two ways in future work: 1. analyze real codebases from production quality games and 2. develop a more complex reference game and ensure its completeness. With

1. real codebases the limitation is that comparative analysis may be impossible because a particular game typically exists as a single implementation only. Analyzing those could still provide useful insight, and might be used to validate the Object-Points analysis in the area of networked games by comparing the source code based complexity measures and the real recorded development effort (the budget that was spent for the programming). To achieve 2. we should carefully specify reference games that cover all areas of networked gaming, but remain small enough to implement within a realistic timeframe (2-4 weeks of programming, reusing common assets) and for qualitative analysis by reading the alternative source codes. Existing canonical implementations may give a starting point, as for example several commercial networked 3d first-person shooter (FPS) games have been open sourced (Quake, Cube2), and at least one highlevel platform already features a FPS as a tutorial (Torque3d).

VI. Conclusions

TODO

(We are happy and curious about using this tool for many kinds of comparisons: longitudal studies of a single API over time, comparisons of e.g. networking stacks when using different protocols for similar functionality, ... or?)

Similarities and differences of using a platform as ready made client software, on which just run scripts, vs. libraries to create own applications, are interesting to study more. Same software components (libraries, modules etc) can be used in both configurations -- what is more suitable may well depend on the particular case.

(XXX Q: where does complexity lurk? should we consider the leaks here? does Onion have something to handle them? at least had the Attribute setting exception in the java server XXX)

References

[api-matters] Michi Henning, API Design Matters, Communications of the ACM Vol. 52 No. 5 http://cacm.acm.org/magazines/2009/5/24646-api-design-matters/fulltext

[cmu-api_failures] Marcelo Cataldo1, Cleidson R.B. de Souza2 (2011). The Impact of API Complexity on Failures: An Empirical Analysis of Proprietary and Open Source Software Systems. http://reports-archive.adm.cs.cmu.edu/anon/isr2011/CMU-ISR-11-106.pdf

[api-complexity-analysis] Comparing Complexity of API Designs: An Exploratory Experiment on DSL-based Framework Integration. http://www.sba-research.org/wp-content/uploads/publications/gpce11.pdf

[pong-ping] High and Low Ping and the Game of Pong. http://www.cs.umu.se/~greger/pong.pdf

[sirikata-scripting] Bhupesh Chandra, Ewen Cheslack-Postava, Behram F. T. Mistree, Philip Levis, and David Gay. "Emerson: Scripting for Federated Virtual Worlds", Proceedings of the 15th International Conference on Computer Games: AI, Animation, Mobile, Interactive Multimedia, Educational & Serious Games (CGAMES 2010 USA). http://sing.stanford.edu/pubs/cgames10.pdf

[henrich97repositorybased] Andreas Henrich, Repository Based Software Cost Estimation, DEXA'97

[conceptmaps]   Jens Gerken, Hans-Christian Jetter, Michael Z llner, Martin Mader, and Harald Reiterer. The concept maps method as a tool to evaluate the usability of apis, May 2011. CHI 2011, May 7–12, 2011, Vancouver, BC, Canada. http://hci.uni-konstanz.de/downloads/CHI2011_concept_maps__publisher_ready.pdf

[overview]   Michael Barth, API Evaluation -- An overview of API evaluation techniques. http://dev.roleplaytalk.net/files/publications/api-evaluation.pdf

[middleware]   T. Hsiao and S. Yuan, "Practical Middleware for Massively Multiplayer Online Games," IEEE Internet Computing, vol. 9, 2005, pp. 47-54.

[apipeerreview]   Farooq, Umer and Welicki, Leon and Zirkler, Dieter, API usability peer reviews: a method for evaluating the usability of application programming interfaces, Proceedings of the 28th international conference on Human factors in computing systems, CHI '10